


Beyond Static Benchmarks: A Post-Leaderboard Evaluation Paradigm for Generative Music and Creative AI

Alexander Lerch
Music Informatics Group, Georgia Institute of Technology



TL;DR

A community-governed, multi-dimensional evaluation ecosystem for music AI replacing static leaderboards with transparent, reproducible *benchmark profiles*.

Problem: Evaluation Crisis

- Generative music evaluation is **methodologically fragmented**, relies on **ambiguous validation objectives**, and builds on **systematically misaligned metrics** [1].
- **Validity failure:** metrics poorly capture musical qualities.
 - **Leaderboard failure:** single scores encourage optimization artifacts and obscure musical multi-dimensionality.
 - **Comparability failure:** different evaluation pipelines yield incomparable results.
 - **Reproducibility failure:** proprietary data and tools limit verification.
 - **Governance failure:** no shared standards or process for evolution.

Core Ideas

- **Community governance:** community-maintained evolving benchmarks and evaluation standards.
- **Multi-dimensional assessment:** capture many *relevant dimensions* instead one single ambiguous score.
- **Interpretability:** utilize *interpretable metrics* and features.
- **Use-case adaptability:** tailor curated *reference benchmark profiles* (RBPs) to specific musical contexts.
- **Modular, extensible architecture:** new metrics, features, and data can be *integrated without redesign*.
- **Reproducibility & transparency:** versioned, documented *open source* software with *auditable logging*.

Governance




Figure: Proposed governance structure

Beyond Music

- modular architecture $\Rightarrow$ easily transferable to other domains of generative AI by **replacing feature representations**; candidates are
 - Visual art
 - Text generation
 - Game content generation
 - Video generation

Framework Architecture




Figure: Proposed framework architecture for the evaluation of generative music.

- **Reference Benchmark Profiles (RBPs):** versioned evaluation configurations specifying
 - reference data
 - feature representations
 - distance metrics
 - reporting format
- **Standardized reference data:** publicly available, documented datasets as controllable open references.
- **Standardized feature representations:** for musical qualities, audio fidelity, controllability, and originality.
- **Standardized distance metrics:**
 - distributional distance metrics (Wasserstein, KL, MMD, etc.)
 - traditional evaluation metrics (accuracy, AUC, F1, etc.) for classification tasks
- **Multi-dimensional results & visualization:**




Figure: Mockup of aggregated result presentation.

- **Logging:** auditable logs (machine & human readable formats) with RBP, version info, test data, etc.

Limitations

- Missing **subjective evaluation protocols**
- **Black-box evaluation:**
 - no evaluation of the underlying model architecture/capabilities
 - no distinction between human contributions and fully machine-generated output
- **Focus on output qualities:**
 - no usability or user experience evaluation
 - no evaluation other characteristics, such as explainability, bias, resource usage, ethical use of data, etc.

References

Alexander Lerch, Claire Arthur, Nick Bryan-Kinns, Corey Ford, Qianyi Sun, and Ashvala Vinay. Survey on the Evaluation of Generative Models in Music. *ACM Computing Surveys*, 58(4):99:1–99:36, October 2025.